

# THE TANK — PITCH REPORT

Agently • Generated 04/08/2026

## Startup Overview

Startup Name: Agently

Funding Ask: ₹150 Lakhs for 8% equity

**Problem Statement:** Busy professionals, founders, freelancers, and students lose significant time managing emails, tasks, calendars, documents, and follow-ups across disconnected apps. Existing AI tools can answer questions, but they do not reliably remember personal context or complete real work safely on a user's behalf.

**Solution:** Agently is a personal AI agent that acts like a digital chief of staff. It connects with a user's everyday tools, learns their preferences and priorities, and helps organize information, plan work, draft communications, manage follow-ups, and complete approved multi-step tasks. Users stay in control through clear permissions and approval before important actions.

**Revenue Model:** We will use a freemium subscription model. A free tier provides limited AI assistance; paid individual plans unlock more agent tasks, integrations, memory, and automation. Over time, we will add premium plans for power users and team plans for small businesses.

**Target Audience:** Our first customers are AI-savvy Indian professionals aged 22–45—especially founders, freelancers, consultants, creators, and knowledge workers—who use Gmail, WhatsApp, calendars, documents, and task apps daily and are willing to pay to save time and reduce mental overload.

---

## Judge Q&A

Victoria Vance asked: "You're targeting Our first customers are AI-savvy Indian professionals aged 22–45—especially founders, freelancers, consultants, creators, and knowledge workers—who use Gmail, WhatsApp, calendars, documents, and task apps daily and are willing to pay to save time and reduce mental overload. — what's your realistic total addressable market, and how does Agently actually capture it at scale?"

Answer: That is exactly why we are not treating "all Indian professionals" as our market.

Our realistic initial serviceable market is roughly 3–5 million digitally intensive Indian knowledge workers—founders, freelancers, consultants, creators, and early AI adopters in major cities—who already pay for productivity software. At an average price of about ₹499 per month, that represents a ₹160 crore annual opportunity.

We do not need to win the entire market. Our first three-year target is 100,000 paying users, which is only around 2–3% of that initial segment and represents roughly ₹160 crore in revenue.

We will capture it through a focused wedge: helping users manage follow-ups, tasks, and communication across the tools they already use. We will acquire early users through founder and freelancer communities, creator partnerships, referral loops, and a free tier that demonstrates value before payment. Once users

trust Agently with a high-frequency workflow, we expand into more tasks, integrations, and paid features. That creates retention and makes the product harder to replace than a one-off AI chatbot.

Reaction: "That's a genuinely scalable answer. I'm leaning in."

Victoria Vance asked: "If a well-funded competitor entered Our first customers are AI-savvy Indian professionals aged 22–45—especially founders, freelancers, consultants, creators, and knowledge workers—who use Gmail, WhatsApp, calendars, documents, and task apps daily and are willing to pay to save time and reduce mental overload.'s space tomorrow, what stops them from crushing you in six months?"

Answer: A well-funded competitor can copy features; they cannot instantly copy user trust, workflow depth, and a focused distribution engine.

Our defensibility is built around becoming the trusted execution layer for a narrow user segment first. Agently will develop a personalized preference and workflow layer—how a user communicates, prioritizes tasks, handles follow-ups, and approves actions. That context compounds with daily use, making the agent more useful over time and creating switching costs.

We will also focus on India-specific needs: WhatsApp-first workflows, local pricing, mobile usage, and the way Indian founders and freelancers manage work across fragmented tools. Large global competitors often build broad products; we will win by solving a few high-frequency workflows exceptionally well.

Finally, our early advantage comes from speed and distribution: embedding in founder, freelancer, and creator communities, earning referrals through visible time savings, and using real user feedback to iterate weekly. If a competitor enters, our goal is not to outspend them—it is to already own the customer relationship, the workflow data, and the habit of daily use.

Reaction: "You just made the market case better than most founders I see."

Marcus Kade asked: "Walk me through the first 90 days after 150 Lakhs f actually ships, and what's the single biggest risk to hitting it?"

Answer: In the first 90 days, we will use the capital to prove one thing: that users will trust Agently to complete a high-frequency workflow and return every week.

**\*\*Days 1–30:\*\*** Build a focused MVP for Gmail, calendar, and task follow-ups. Recruit 30–50 design partners from founders, freelancers, and consultants; interview them weekly; and define the approval and privacy system so the agent never takes important actions without user confirmation.

**\*\*Days 31–60:\*\*** Launch a private beta. The agent will identify follow-ups, turn messages into tasks, draft replies, schedule reminders, and provide a daily priority brief. We will measure activation, weekly retention, completed tasks, and how often users approve suggested actions.

**\*\*Days 61–90:\*\*** Improve the workflows based on usage, add one high-demand integration—likely WhatsApp or a task-management tool—and begin converting the most active users to a paid pilot. Our target is 100 active beta users, at least 40% four-week retention, and 15–20 users willing to pay.

The single biggest risk is not building the technology; it is earning trust. People will only delegate meaningful work if the agent is accurate, transparent, and easy to control. We reduce that risk by starting with low-risk, reviewable actions, requiring approval before sending or changing anything important, and measuring trust through repeat usage and approval rates before expanding automation.

Reaction: "I'd want you running my engineering team with an answer like that."

Marcus Kade asked: "What's the hardest part of actually building Agently is a personal AI agent that acts like a digital chief of staff. It connects with a user's everyday tools, learns their preferences and priorities, and helps organize information, plan work, draft communications, manage follow-ups, and complete approved multi-step tasks. Users stay in control through clear permissions and approval before important actions. that most outsiders underestimate?"

Answer: The hardest part is not the AI model or the integrations. It is building reliable judgment at the boundary between "helpful suggestion" and "taking action."

A personal agent has to understand incomplete context, conflicting priorities, personal preferences, and the consequences of small mistakes. Drafting an email is easy; knowing whether it should be sent, when to follow up, what can be delegated, and when to interrupt the user requires trust, memory, permissions, and strong safeguards.

Most outsiders underestimate how quickly one incorrect action can destroy confidence. If Agently misses an important follow-up, sends the wrong message, or surfaces private information improperly, users will stop delegating work.

That is why we are designing it as a trust product first: narrow workflows, clear explanations, user approval for consequential actions, reversible actions wherever possible, and gradual automation only after the agent has demonstrated accuracy. The real product challenge is making the agent dependable enough that users are comfortable relying on it every day.

Reaction: "Now THAT sounds like someone who's actually shipped before."

Priya Anand asked: "Why would someone actually switch to Agently over whatever they use today to deal with busy professionals, founders, freelancers, and students lose significant time managing emails, tasks, calendars, documents, and follow-ups across disconnected apps. existing ai tools can answer questions, but they do not reliably remember personal context or complete real work safely on a user's behalf.?"

Answer: People will not switch just for another AI chat interface. They will switch if Agently saves them meaningful time without forcing them to abandon the tools they already use.

Today, users manage work across Gmail, WhatsApp, calendars, documents, and task apps. They must remember what needs attention, manually copy information between tools, and repeatedly explain context to AI assistants. Agently works across those existing tools and turns fragmented information into action: it identifies follow-ups, creates prioritized tasks, prepares drafts, and gives users one clear view of what matters.

The key difference is persistence and execution. Most AI tools answer a prompt and forget the context; Agently remembers approved preferences and priorities, proactively helps with recurring work, and

completes multi-step tasks only after the user approves. That makes it feel less like a chatbot and more like a reliable digital chief of staff.

We will earn switching through a simple promise: connect your existing tools, let Agently handle one painful workflow well, and show measurable time saved in the first week.

Reaction: "I would genuinely tell my friends about that."

Priya Anand asked: "What happens the day I stop using Agently — do I even notice it's gone?"

Answer: Yes—and that is the standard we are designing for.

The day a user stops using Agently, they should feel the return of the mental load: missed follow-ups, scattered tasks, repeated planning, and the need to manually check multiple apps to know what matters. Agently's value is not another place to work; it is the invisible coordination work that keeps work moving.

We will measure this through behavior, not slogans: how many follow-ups are caught, tasks are completed, hours are saved, and users return weekly. If users do not notice its absence, then we have not earned a place in their workflow.

Our goal is for Agently to become a trusted daily layer between a user and their existing tools—quietly preparing priorities, keeping commitments on track, and reducing the number of things they have to remember themselves.

Reaction: "I would genuinely tell my friends about that."

Alexander Cole asked: "Your model is We will use a freemium subscription model. A free tier provides limited AI assistance; paid individual plans unlock more agent tasks, integrations, memory, and automation. Over time, we will add premium plans for power users and team plans for small businesses. — walk me through the unit economics: what does it actually cost you to acquire and serve one customer?"

Answer: At launch, these are planning assumptions—not claimed results—and we will validate them through the beta.

Our paid plan is expected to be around  $\$1499$  per month, or roughly  $\$14,999$

For one paid customer, we estimate monthly serving costs of approximately

- AI model and agent-processing costs:  $\$170 - \$1100$
- Cloud, integrations, storage, and monitoring:  $\$130 - \$145$
- Payment processing and customer support:  $\$120 - \$125$

That gives us a gross margin of roughly 65–75% at launch, with room to improve as usage patterns become clearer and model costs decline.

For acquisition, we will initially rely on founder communities, creator partnerships, referrals, and product-led conversion rather than paid advertising. Our target customer-acquisition  $\$1499$  per month, we recover that cost in around two to three months of gr

The important constraint is free-tier cost. We will cap agent tasks and expensive integrations on the free plan, so free users can experience the value without creating unlimited model costs. A healthy target is that one paying customer supports several free users while we convert the most active ones.

If we achieve 12 months of average paid retention, a customer generates  $\text{\text{₹}}14,000$  in gross profit. That creates a sustainable model only remains below the first few months of gross profit—so those are the two metrics we will watch most closely.

Reaction: "Now that is a profitable-looking machine."

Alexander Cole asked: "When do you actually turn a profit, and what happens to that timeline if growth comes in at half your projection?"

Answer: We expect to reach operating profitability at roughly 750–900 paying users, targeted around month 18.

At  $\text{\text{₹}}1499$  per month, 800 paying users produce about  $\text{\text{₹}}4$  lakh in monthly revenue. At  $\text{\text{₹}}11.2$  lakh in monthly costs of roughly  $\text{\text{₹}}1.2$  lakh, we have around  $\text{\text{₹}}2.8$  lakh in contribution margin. At an operating cost of approximately  $\text{\text{₹}}2.5\text{--}2.8$  lakh for a small product, engineering team.

The  $\text{\text{₹}}50$  lakh raise is designed to give us about 18 months of runway while we do not hire ahead of retention or revenue; the early team stays deliberately lean.

If growth is half our projection, profitability moves out to roughly month 28–30. In that case, we do not simply spend faster and hope: we slow hiring, reduce paid acquisition, focus on the highest-retention customer segment, and extend runway. The decision point is month 6: if retention or conversion is weak, we narrow the product to the workflow users value most before scaling spend.

The key is that our break-even point is relatively low because this is subscription software. We do not need millions of users to become sustainable—but we do need strong retention and disciplined operating costs.

Reaction: "Now that is a profitable-looking machine."

---

## Scorecard

Market Potential: 75/100

Innovation: 80/100

Business Model: 81/100

Execution: 68/100

Investment Worthiness: 76/100

Overall Score: 76/100

---

## Verdict: They're In

Funding / Offer Amount: 150 L

Implied Valuation: 12.27 Cr

Equity Requested: 22.0%

The panel is convinced. Agently's strongest signal was business model at 81/100. Combined with solid answers under pressure, this cleared the bar for a check.

The clearest gap was execution, scoring 68/100 — that's the first thing to fix. Still, the overall picture at 76/100 is fundable today.

Expect the sharks to want board visibility and a clear path to the next milestone before the next round.